


Section 5:

1. Movie Review Data

- <http://www.cs.cornell.edu/people/pabo/movie-review-data/>
 - Download [polarity dataset v2.0](#) (3.0Mb)
 - 1000 positive and 1000 negative processed reviews
- *Extract*



2. Loading data

```
In [4]: from sklearn import datasets

# point to your data location
data_directory = 'review_polarity'
movie_sentiment_data = datasets.load_files(data_directory, shuffle=True)
```

```
In [6]: movie_sentiment_data.data
```

```
s about flintstones , family set in fictitious stone age " town " of bedrock , whose members enjoy the lifestyle of 1950s m
iddle class america . \nfred flintstone ( john goodman ) works in the quarry , and one day he helps his best friend and ne
ighbour barney rubble ( rick moranis ) and his wife betty ( rosie o\donnell ) to adopt a baby . \nto return the favour ,
barney switches his results of aptitude test with fred , and , based on that , fred gets well-paid job in management . \n
but , of course , this is just a sham - corrupt official cliff vandercafe ( kyle macclachlan ) and his sultry secretary sh
aron stone ( halle berry ) need a scapegoat for their embezzlement scheme . \nin the meantime , fred\ 's wife wilma ( eliz
abeth perkins ) must face her mother ( elizabeth taylor ) who can\ 't stand fred . \non the superficial level , the flints
tones did the excellent job in bringing the animated series to life . \ncomputer effects are flawless , and the costumes
, settings and other details are authentic for all the fans of the show . \nunfortunately , problems with this film start
with inadequate casting - rick moranis is too thin for the role of barney , while the cartoon betty used to be much skinn
ier than rosie o\donnell . \nbut the greatest problem of all is plot , or to be precise , the lack of plot . \nsome thir
ty six screenwriters made sure that the plot of the film is lame , original characters one-dimensional , and many element
s of the story , like embezzlement and inter-office politics , totally incomprehensible for the little children , the mai
n targeted audience of this film . \nresult is almost unwatchable mess , occasionally saved mostly by excellent acting (
elizabeth perkins was right on mark as wilma ) and one of the classic example of mortal hollywood disease known as " high
concept " . \nafter great hype , movie quickly sank into oblivion and the fans of the show returned to the animated versi
on . \nall in all , film isn\ 't that bad , but only the hard core fans and nostalgics can find more than guilty pleasure
in it . \n' ,
...]
```

```
In [7]: movie_sentiment_data filenames
```

```
Out[7]: array(['review_polarity\\neg\\cv405_21868.txt',
               'review_polarity\\pos\\cv190_27052.txt',
               'review_polarity\\pos\\cv132_5618.txt', ...,
               'review_polarity\\pos\\cv653_19583.txt',
               'review_polarity\\neg\\cv559_0057.txt',
               'review_polarity\\neg\\cv684_12727.txt'], dtype='<U35')
```

```
In [10]: movie_sentiment_data.target_names
```

```
Out[10]: ['neg', 'pos']
```

3. Preprocessing the dataset

```
In [14]: from sklearn import feature_extraction
import nltk
```

```
In [15]: def extract_features(corpus):
'''Extract TF-IDF features from corpus'''
# vectorize means we turn non-numerical data into an array of numbers
count_vectorizer = feature_extraction.text.CountVectorizer(
    lowercase=True, # for demonstration, True by default
    tokenizer=nltk.word_tokenize, # use the NLTK tokenizer
    stop_words='english', # remove stop words
    min_df=1 # minimum document frequency, i.e. the word must appear more than once.
)
processed_corpus = count_vectorizer.fit_transform(corpus)
processed_corpus = feature_extraction.text.TfidfTransformer().fit_transform(
    processed_corpus)

return processed_corpus
```

```
In [16]: movie_tfidf = extract_features(movie_sentiment_data.data)
```

```
In [17]: movie_tfidf
```

```
Out[17]: <2000x46156 sparse matrix of type '<class 'numpy.float64'>'
with 506362 stored elements in Compressed Sparse Row format>
```

4. Logistic regression model

```
In [21]: # Split train, test sets
from sklearn.model_selection import train_test_split
```

```
In [22]: X_train, X_test, y_train, y_test = train_test_split(
    movie_tfidf, movie_sentiment_data.target, test_size=0.30, random_state=42)
```

```
In [23]: # Call logistic regression model
from sklearn.linear_model import LogisticRegression
```

```
In [24]: model = LogisticRegression()
model.fit(X_train, y_train)

print('Model performance: {}'.format(model.score(X_test, y_test)))

Model performance: 0.7983333333333333
```

```
In [28]: p_pred = model.predict(X_test)
```

```
In [33]: from sklearn.metrics import plot_confusion_matrix
import matplotlib.pyplot as plt
```

```
In [34]: classifier = model
class_names = movie_sentiment_data.target_names

# Plot non-normalized confusion matrix
titles_options = [("Confusion matrix, without normalization", None),
                  ("Normalized confusion matrix", 'true')]
for title, normalize in titles_options:
    disp = plot_confusion_matrix(classifier, X_test, y_test,
                                display_labels=class_names,
                                cmap=plt.cm.Blues,
                                normalize=normalize)

    disp.ax_.set_title(title)

    print(title)
    print(disp.confusion_matrix)

plt.show()
```

5. Confusion matrix

Confusion matrix, without normalization

```
[[228  58]
 [ 63 251]]
```

Normalized confusion matrix

```
[[0.7972028  0.2027972 ]
 [0.20063694 0.79936306]]
```

